

CSE 753: Advanced Reinforcement Learning

Contents

1	Policy Gradient Methods	2
2	Actor Critic Method	23

1 Policy Gradient Methods

1.1 Generalized Policy Iteration

Generalized Policy Iteration — the core idea behind how an AI agent learns to make better decisions over time.

The Big Idea

The agent follows a simple two-step loop, repeated over and over:

1. **Collect experience** — Run the current policy (the agent’s current strategy, called π) out in the world and record what happened. These recorded sequences of actions and outcomes are called **trajectories** — basically a replay of what the agent did from start to finish.
2. **Improve the policy** — Look at all those trajectories. Which ones led to good outcomes? Which led to bad ones? Use that information to update the policy so it’s slightly smarter than before.

Then repeat. That’s it. The whole loop.

What “Initialize” Means

Before the loop starts, you need to give the agent some starting policy. You have three options:

- **Random** — just pick actions randomly at first (agent knows nothing)
- **Imitation learning** — copy how an expert does it
- **Heuristics** — use simple hand-coded rules as a starting point

How It Improves Over Time

The diagram on the right shows three snapshots — Iteration 1, 50, and 100. Think of each curved line as one trajectory (one attempt the agent made from start to finish).

- Red X = bad outcome (low reward)
- Yellow check = okay outcome
- Green check = good outcome

Stage	What’s happening
Early (Iter 1)	Most attempts fail. The agent is basically guessing. Lots of exploration.
Mid (Iter 50)	Getting better. Some trajectories succeed. Fewer failures.
Late (Iter 100)	Almost all trajectories lead to good outcomes. The policy has converged.

The Key Takeaway

The agent doesn’t need to know the rules of the world in advance. It just tries things, sees what worked, updates itself, and tries again. Over hundreds of iterations, this simple loop produces a very good policy.

1.2 Online Learning

Is this online learning only?

No. The slide specifically describes a **batch** update style — it says “collect a **batch** of data (trajectories)”, then update. That’s not online learning.

Here’s the distinction:

- **Online learning** — update the policy after **every single step** (or every transition). You don’t wait.
- **Batch learning** — collect a **full set of trajectories** first, then update the policy once using all of them. That’s what the slide shows.
- **Mini-batch** — a middle ground; collect a fixed number of episodes, update, repeat.

GPI as a **general framework** doesn’t lock you into one style. It just says “evaluate → improve → repeat.” How often you update is an implementation choice. That’s why both REINFORCE (batch, full episodes) and Actor-Critic methods (online or near-online, step-by-step) are both considered instances of GPI.

What happens in online learning?

In online learning, the agent updates the policy **after every single step** — no waiting, no batching.

The loop is:

1. Take one action in the environment
2. See the result (reward + next state)
3. Immediately update the policy
4. Take the next action using the already-updated policy

So the policy is changing **continuously** as the agent moves through the world. By the time the episode ends, the policy has already been updated dozens or hundreds of times.

The tradeoff: updates are fast and frequent, but each update is noisy because it’s based on just one data point instead of a whole batch of experience.

In online learning agent interacts with real environment directly. Right?

Not necessarily. “Online learning” just means **updating after every step** — it says nothing about whether the environment is real or simulated.

- In a **real environment** — yes, the agent takes a real action, gets a real reward, updates immediately.
- In a **simulated environment** — the agent still updates after every step, it’s still online learning. The environment just happens to be a simulator.

The word **online** refers to the **frequency of updates**, not the **type of environment**.

The distinction you might be thinking of is **Online (real environment)** vs. **Offline / Batch RL (replay buffer or pre-collected dataset)** — here “offline” means the

agent learns from a fixed dataset with no new interactions at all. But that’s a different axis than simulated vs. real.

1.3 Policy Without Q Values

Page 3 is titled “**Policy without Q values**” and introduces a new way of thinking about policies.

The Setup

So far in the course, GPI-based algorithms figured out a good policy by first computing $Q(s, a)$ — the expected reward of taking action a in state s . Once you know Q values well enough, you derive the policy from them (e.g., pick the action with the highest Q value). The slide says: what if we skip Q values entirely and just **directly build the policy itself**?

The New Idea: Parameterized Policy

Instead of storing Q values and deriving the policy from them, we represent the policy as a **mathematical function with adjustable knobs**. Those knobs are called **parameters**. We then tune the knobs directly until the policy gets good. This is called **parameterization** — you describe the policy using a set of numbers (the parameters), and learning means finding the right numbers.

The Equation

$$\pi(a \mid s, \vec{\theta}) = \Pr\{A_t = a \mid S_t = s, \vec{\theta}_t = \vec{\theta}\}$$

Symbol glossary:

Symbol	Meaning
--------	---------

π	The policy — the agent’s decision-making rule
a	A specific action the agent could take
s	The current state the agent is in
$\vec{\theta}$	The parameter vector — a list of numbers that defines the shape of the policy
d'	The number of parameters in $\vec{\theta}$ (its size/dimension)
$\Pr\{\cdot\}$	Probability of something happening
A_t	The action actually taken at time step t
S_t	The state the agent was in at time step t
t	A specific moment in time during the episode

The policy $\pi(a \mid s, \vec{\theta})$ gives you the **probability of choosing action** a , given that you’re in state s and your policy’s parameters are currently set to $\vec{\theta}$. The policy is no longer a lookup table — it’s a function. You plug in the state and the parameters, and it spits out a probability for each possible action. Learning now means **adjusting** $\vec{\theta}$ to make good actions more probable and bad actions less probable.

1.4 Parameterizing Policy

Page 4 is titled “**Parameterizing Policy**” and answers a practical question: if the policy is just a function with parameters, what rules must it follow, and what’s a concrete

example of such a function?

Two Rules Every Policy Must Obey

Since $\pi(a \mid s, \theta)$ outputs a **probability**, it has to follow the basic laws of probability. The slide gives two constraints.

Rule 1

$$\pi(a \mid s, \theta) \geq 0 \quad \text{for all } a \in \mathcal{A} \text{ and } s \in \mathcal{S}$$

Symbol glossary:

Symbol	Meaning
--------	---------

$\pi(a \mid s, \theta)$	Probability of taking action a in state s with parameters θ
$\mathcal{A}$	The set of all possible actions
$\mathcal{S}$	The set of all possible states

Every probability the policy outputs must be zero or positive. You can't have a negative probability — that makes no sense. This is just the most basic rule of probability.

Rule 2

$$\sum_{a \in \mathcal{A}} \pi(a \mid s, \theta) = 1 \quad \text{for all } s \in \mathcal{S}$$

Symbol glossary:

Symbol	Meaning
--------	---------

$\sum_{a \in \mathcal{A}}$	Add up over every possible action in $\mathcal{A}$
----------------------------	--

If you add up the probabilities of all possible actions in any given state, the total must equal 1. The agent has to do *something* — so all options together must account for 100% of the probability.

The Softmax Function — A Concrete Example

$$\pi(a \mid s, \theta) = \frac{e^{h(s,a,\theta)}}{\sum_{b \in \mathcal{A}} e^{h(s,b,\theta)}}$$

Symbol glossary:

Symbol	Meaning
--------	---------

$h(s, a, \theta)$	The preference score for action a in state s — a raw number produced by the parameter
$e^{(\cdot)}$	The exponential function — turns any number into a positive value, amplifies differences
b	A dummy variable representing each action in the sum (used to loop over all actions)
$\sum_{b \in \mathcal{A}} e^{h(s,b,\theta)}$	Total of exponential scores across all actions — used to normalize

First, compute a preference score $h(s, a, \theta)$ for every action. Then take e to the power of each score to make them all positive. Finally, divide each action's score by the total of all scores, so everything adds up to 1. The result is a valid probability for each action. Actions with higher preference scores get higher probabilities, but no action is ever completely ignored (its probability never hits exactly zero).

The Greedy Behavior Note

If the preference $h(s, a, \theta)$ of a particular action grows very large, the policy tends to act greedily. This means: if one action's score gets much bigger than all others, the softmax will assign it nearly 100% probability and ignore the rest. The policy stops exploring and just always picks that one action — it becomes **greedy**.

1.5 Preference Score and Feature Vectors

Who defines the preference score? How is it defined?

The parameters θ define it — specifically, $h(s, a, \theta)$ is a function that you design, and θ are the learnable weights inside that function.

How it's defined in practice

The most common choice is a **linear function**:

$$h(s, a, \theta) = \theta^\top x(s, a)$$

where $x(s, a)$ is a feature vector you hand-craft to describe the state-action pair, and θ is a weight vector the algorithm learns. The dot product gives a single score.

A more powerful choice is a **neural network** — you feed the state s in, and the network outputs one score per action. The weights of the network are θ , and they get updated through gradient descent.

So who decides the structure?

- **You (the designer)** choose the form of h — linear, neural network, etc.
- **The learning algorithm** (e.g., REINFORCE) then adjusts θ automatically over training to make good actions score higher and bad actions score lower.

The preference score has no inherent meaning on its own — it's just a raw number. What matters is the **relative difference** between scores across actions, because softmax only cares about which score is bigger, not the absolute values.

“ $x(s, a)$ is a feature vector you hand-craft”. What do you mean by hand-craft?

It means **you manually decide** what information to extract from the state-action pair and feed into the function.

Example — a robot navigating a grid

Say the state s is the robot's position, and action a is “move right.” You might hand-craft the feature vector as:

$$x(s, a) = \begin{bmatrix} \text{distance to goal} \\ \text{distance to wall} \\ 1 \text{ if action} = \text{right, else } 0 \\ 1 \text{ if action leads to open space, else } 0 \end{bmatrix}$$

You — the human — decided that these four pieces of information matter. You wrote that logic manually. That’s what “hand-crafted” means.

Why is this a problem?

- You might miss important features
- You might include useless ones
- For complex problems (images, text, games), it’s nearly impossible to hand-craft good features

That’s exactly why neural networks replaced hand-crafted features

A neural network takes the raw state (e.g., raw pixels) and **learns its own internal features automatically** from experience. You don’t decide what matters — the network figures it out. The weights θ inside the network serve the same role as the hand-crafted $x(s, a)$, but nothing is manually designed. So hand-crafting was the old way. Neural networks are the modern way.

1.6 The RL Objective

Page 6 is titled “**The RL Objective**” — it formally states what the agent is actually trying to maximize, and how to estimate it in practice.

The Big Goal

The agent wants to find the best possible parameters θ for its policy. “Best” means: the parameters that produce the highest total reward on average across all possible trajectories.

Equation 1 — Finding the Best Parameters

$$\vec{\theta}^* = \arg \max_{\theta} \mathbb{E}_{\tau \sim p_{\theta}(\tau)} [r(\tau)]$$

Symbol glossary:

Symbol	Meaning
$\vec{\theta}^*$	The optimal (best) parameter vector — what we’re searching for
$\arg \max_{\theta}$	“Give me the value of θ that makes the thing inside as large as possible”
τ	A trajectory — one full sequence of states and actions from start to finish
$p_{\theta}(\tau)$	The probability of trajectory τ happening under policy θ
$\mathbb{E}_{\tau \sim p_{\theta}(\tau)}[\cdot]$	The average value of something, taken over all trajectories the policy generates
$r(\tau)$	The total reward collected along trajectory τ

We want the parameters θ^* that make the **average total reward across all trajectories** as high as possible. We're not optimizing for one lucky run — we want good performance on average.

Equation 2 — Expanding the Reward

$$= \arg \max_{\theta} \mathbb{E}_{\tau \sim p_{\theta}(\tau)} \left[\underbrace{\sum_t r(s_t, a_t)}_{J(\theta)} \right]$$

Symbol glossary:

Symbol Meaning

$\sum_t$ Add up over every time step t in the trajectory

$r(s_t, a_t)$ The reward received at time step t , for being in state s_t and taking action a_t

$J(\theta)$ The objective function — the single number we want to maximize; our score for how good

The total reward of a trajectory $r(\tau)$ is just the sum of all the small rewards collected at each individual step. $J(\theta)$ is the name we give to the expected value of that sum — it's our score for how good a given θ is.

Equation 3 — How to Actually Compute It

$$J(\theta) \approx \frac{1}{N} \sum_i \sum_t r(s_{i,t}, a_{i,t})$$

Symbol glossary:

Symbol Meaning

$\approx$ Approximately equal to

N Number of trajectories (sample runs) collected

$\sum_i$ Add up over each trajectory i (each run)

$\sum_t$ Add up over each time step within that trajectory

$r(s_{i,t}, a_{i,t})$ Reward at time step t in trajectory i

$\frac{1}{N}$ Divide by number of trajectories to get the average

You can't compute the true expectation $J(\theta)$ exactly — that would require running infinitely many trajectories. Instead, you run N trajectories, add up all the rewards collected across all time steps in all trajectories, and divide by N to get the average. This average is a good enough approximation of the true $J(\theta)$.

The Diagram on the Right

This is exactly what the approximation looks like in practice. The dot on the left is the starting state. Each curved arrow is one trajectory — one attempt by the agent. Some end in red X (bad outcome, low reward), some in yellow or green checkmark (good or great outcome). You run N of these, average the total rewards, and that gives you your estimate of $J(\theta)$.

1.7 Pages 7–8: Deriving the Gradient

Page 7 — Gradient of the RL Objective

Quick recap. Remember from page 6: the agent’s goal is to maximize $J(\theta)$, the average total reward it gets by following its policy. To improve the policy, we need to compute the **gradient** of $J(\theta)$ — basically, the direction in which to nudge the policy knobs (θ) to get more reward.

Equation 1 — Rewriting the objective as an integral.

$$J(\theta) = E_{\tau \sim p_\theta(\tau)} \left[\sum_t r(s_t, a_t) \right] = \int p_\theta(\tau) r(\tau) d\tau$$

Symbol	Meaning
--------	---------

$J(\theta)$	The score of the policy — total expected reward
-------------	---

θ	The policy’s knobs (parameters)
----------	---------------------------------

τ	A trajectory — a full playthrough: $s_1, a_1, s_2, a_2, \dots$
--------	--

$p_\theta(\tau)$	Probability of that trajectory happening under policy θ
------------------	--

$r(\tau)$	Total reward collected along that trajectory
-----------	--

$\int \dots d\tau$	“Sum over all possible trajectories”
--------------------	--------------------------------------

This is just rewriting the expected reward in math notation. Instead of saying “average reward over many runs,” we write it as an integral — summing $p_\theta(\tau) \cdot r(\tau)$ over every possible trajectory. The integral is mathematically equivalent to the expectation, which sets us up for taking the gradient.

Equation 2 — Taking the gradient (the log-derivative trick).

$$\nabla_\theta J(\theta) = \int \nabla_\theta p_\theta(\tau) r(\tau) d\tau = \int p_\theta(\tau) \nabla_\theta \log p_\theta(\tau) r(\tau) d\tau = E_{\tau \sim p_\theta} [\nabla_\theta \log p_\theta(\tau) r(\tau)]$$

Symbol	Meaning
--------	---------

∇_θ	“How much does this change when we tweak θ ?” (the gradient)
-----------------	---

$\nabla_\theta p_\theta(\tau)$	Gradient of the trajectory probability — hard to work with directly
--------------------------------	---

$\nabla_\theta \log p_\theta(\tau)$	Gradient of the <i>log</i> of the trajectory probability — much easier
-------------------------------------	--

We want the gradient of $J(\theta)$, but differentiating $p_\theta(\tau)$ directly is messy because it involves the unknown environment dynamics. The **log-derivative trick** says: $p_\theta(\tau) \cdot \nabla_\theta \log p_\theta(\tau) = \nabla_\theta p_\theta(\tau)$. This lets us swap the hard term out and rewrite the whole gradient as an **expectation** — which means we can estimate it just by running the policy and averaging, no integrals needed.

The key identity.

$$p_\theta(\tau) \nabla_\theta \log p_\theta(\tau) = p_\theta(\tau) \cdot \frac{\nabla_\theta p_\theta(\tau)}{p_\theta(\tau)} = \nabla_\theta p_\theta(\tau)$$

This is just the chain rule of calculus: the derivative of $\log(x)$ is $\frac{1}{x}$. Multiplying $p_\theta(\tau)$ on both sides cancels it out, proving the substitution above is valid.

Page 8 — Expanding the Gradient into Something Computable

Now the slide answers: “Okay, but how do we actually **compute** $\nabla_\theta \log p_\theta(\tau)$?”

Equation 3 — What is $p_\theta(\tau)$?

$$p_\theta(\tau) = p(s_1) \prod_{t=1}^T \pi_\theta(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

Symbol	Meaning
$p(s_1)$	Probability of starting in state s_1 (set by the environment, not the agent)
$\prod_{t=1}^T$	Multiply together across all time steps $t = 1$ to T
$\pi_\theta(a_t s_t)$	Policy’s probability of picking action a_t when in state s_t
$p(s_{t+1} s_t, a_t)$	Environment’s probability of moving to s_{t+1} given state and action

A trajectory’s probability is the product of every decision and every world transition along the way. Think of it like the probability of a specific path through a maze — it’s the starting probability times the chance you took each turn times the chance the world responded the way it did, multiplied across every step.

Equation 4 — Expanding $\log p_\theta(\tau)$.

$$\nabla_\theta \left[\log p(s_1) + \sum_t \log \pi_\theta(a_t | s_t) + \log p(s_{t+1} | s_t, a_t) \right]$$

Taking the log of a product turns it into a sum (since $\log(ABC) = \log A + \log B + \log C$). The critical insight is that $\log p(s_1)$ and $\log p(s_{t+1} | s_t, a_t)$ do not contain θ at all — they are controlled by the environment, not the policy. So their gradients are **zero** and they vanish. Only $\log \pi_\theta(a_t | s_t)$ survives.

Equation 5 — The practical gradient formula.

$$\nabla_\theta J(\theta) = E_{\tau \sim p_\theta} \left[\left(\sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t | s_t) \right) \left(\sum_{t=1}^T r(s_t, a_t) \right) \right]$$

Symbol	Meaning
$\sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t s_t)$	How to nudge the policy to make each action more or less likely
$\sum_{t=1}^T r(s_t, a_t)$	Total reward collected during the trajectory

The gradient is the product of two sums — “how much did each action’s probability change when we tweak θ ?” multiplied by “how much total reward did we get?” If a trajectory got high reward, we push θ in the direction that made those actions more

likely. If it got low reward, we push the other way. Crucially, the environment dynamics $p(s_{t+1} \mid s_t, a_t)$ are **completely gone** — we only need the policy and the rewards.

Equation 6 — The sample estimate (how we actually compute it).

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{i,t} \mid s_{i,t}) \right) \left(\sum_{t=1}^T r(s_{i,t}, a_{i,t}) \right)$$

Symbol	Meaning
N	Number of sampled trajectories (rollouts)
i	Index for each sampled trajectory
$\frac{1}{N} \sum_i$	Average over all N rollouts
$a_{i,t}, s_{i,t}$	Action and state at time t in the i -th trajectory

Since we cannot compute an expectation over *all* possible trajectories, we approximate it by running the policy N times, computing the gradient formula for each run, and averaging. This is the same idea as estimating the average of a die roll by rolling it many times — the more rollouts, the better the estimate. This is the actual algorithm you would implement in code.

1.8 Page 9: The Full REINFORCE Algorithm

This page puts everything from pages 7 and 8 together into one complete, runnable algorithm. It has a name: **REINFORCE**, also called the **vanilla policy gradient**. “Vanilla” just means the plain, basic version — no fancy extras yet.

Think of it like a recipe with 4 steps that you repeat over and over until the agent gets good.

Step 1 — Collect experience.

Sample $\{\tau^i\}$ from $\pi_{\theta}(a_t \mid s_t)$

Symbol	Meaning
$\{\tau^i\}$	A collection of trajectories (multiple full playthroughs), indexed by i
$\pi_{\theta}(a_t \mid s_t)$	The current policy — given state s_t , it picks action a_t

Run your current policy in the environment multiple times and record what happened each time — which states you visited, which actions you took, and what rewards you got. This is just the agent “playing the game” several times and taking notes.

Step 2 — Compute the gradient.

$$\nabla_{\theta} J(\theta) \approx \sum_i \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i \mid s_t^i) \right) \left(\sum_{t=1}^T r(s_t^i, a_t^i) \right)$$

Symbol	Meaning
$\nabla_{\theta} J(\theta)$	The gradient — which direction to push the policy knobs
$\sum_i$	Sum (add up) over all collected trajectories
$\nabla_{\theta} \log \pi_{\theta}(a_t^i s_t^i)$	How much action a_t in trajectory i would change if we tweaked θ
$\sum_{t=1}^T r(s_t^i, a_t^i)$	Total reward collected in trajectory i

For each recorded trajectory, multiply “how sensitive was the policy to its own choices” by “how much reward did it get?” Add this up across all trajectories. Trajectories with high reward push the gradient hard in a direction that makes those actions more likely; trajectories with low reward push in the opposite direction. This is the signal that tells us how to improve.

Step 3 — Update the policy.

$$\vec{\theta} = \vec{\theta} + \alpha \nabla_{\theta} J(\theta)$$

Symbol	Meaning
$\vec{\theta}$	The policy’s parameters (knobs) — a vector of numbers
α	Learning rate — controls how big a step to take
$\nabla_{\theta} J(\theta)$	The gradient computed in Step 2

Nudge the policy parameters θ in the direction of the gradient by a small amount controlled by α . Note it is a **plus** sign, not minus — because we are doing **gradient ascent** (climbing toward higher reward), unlike regular gradient descent used in supervised learning which minimizes a loss. A large α means big jumps (risky); a small α means careful, slow improvement.

Step 4 — Repeat until convergence.

Loop back to Step 1 with the newly updated policy. Collect new data, compute a new gradient, update again. Keep going until the policy stops improving noticeably — that is convergence.

The big picture. The diagram on the slide shows this loop: collect experience by running the policy, then improve the policy using that data, then repeat. REINFORCE looks at all collected trajectories, figures out which actions led to good outcomes, and makes those actions more likely next time. It is literally trial and error, formalized into math.

1.9 Page 10: Understanding What the Gradient Actually Does

This page is not introducing new math. It is stepping back and asking: **what does this gradient formula actually mean in plain English?** It is the intuition behind everything derived on pages 7–9.

The gradient (same formula, new lens).

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \right) \left(\sum_{t=1}^T r(s_t^i, a_t^i) \right)$$

Symbol	Meaning
$\frac{1}{N} \sum_i$	Average over all N collected trajectories
$\nabla_{\theta} \log \pi_{\theta}(a_t^i s_t^i)$	Direction to push θ to make action a_t in trajectory i more likely
$\sum_{t=1}^T r(s_t^i, a_t^i)$	Total reward of trajectory i — the weight

The total reward of a trajectory acts as a **multiplier** on the gradient. If a trajectory earned high reward, its gradient contribution gets multiplied by a big positive number — meaning those actions get strongly encouraged. If the reward was negative, the multiplier flips the sign and those actions get discouraged. The gradient is literally a weighted vote from every trajectory.

The three key points from the slide.

1. **Increase the likelihood of actions you took in high-reward trajectories.** If a playthrough went well, assume the actions you took during it were probably good. Make those actions more likely in the future.
2. **Decrease the likelihood of actions you took in negative-reward trajectories.** If a playthrough went badly, assume those actions were probably bad. Make them less likely going forward.
3. **Do more of the good stuff, less of the bad stuff.** That is the whole algorithm in one sentence. The math just makes this idea precise and automatic.

The diagram. The picture shows several trajectories starting from the same point. Some end in failure and one ends in success. After running REINFORCE, the policy will shift probability mass toward whatever sequence of actions led to the successful path, and away from the actions that led to failure. Over many iterations, the successful path becomes more and more likely.

Why this page matters. Everything on pages 7 and 8 was dense derivation. This page is the payoff — it tells you that all that math ultimately boils down to one dead-simple idea: **reward good behavior, punish bad behavior, repeat.** REINFORCE is just that idea, expressed precisely enough for a computer to execute.

1.10 Pages 11–12: A Problem with the Gradient and the Reward-to-Go Fix

Page 11 — Quiz: A Problem With the Gradient

This page gives a concrete worked example to expose a flaw hiding inside the gradient formula. The task: **teach a robot to walk forward.**

The reward is simple:

$$r(s, a) = \text{forward velocity of the robot (negative if it goes backwards)}$$

The agent runs 5 trajectories (τ^1 through τ^5):

Trajectory	What happened	Reward
τ^1	Falls backwards	Negative
τ^2	Falls forward	Small positive
τ^3	Stands still	Zero
τ^4	One small step forward, then falls backwards	Mixed
τ^5	One large step backwards, then small step forwards	Negative overall

The question the slide asks: given these 5 trajectories and the gradient formula from page 9, **what will the gradient actually encourage the policy to do?**

The answer: **it will encourage the robot to fall forward — not to actually take a step forward.**

Why? The flaw explained.

τ^2 (falls forward) gets a positive total reward because at least it was moving in the right direction for a moment. τ^4 (takes a step, then falls backward) might net out to a small or zero total reward because the backward fall cancels the forward step.

So when the gradient formula multiplies each trajectory's actions by its **total reward**, τ^2 gets a stronger positive push than τ^4 . The gradient is essentially saying: “falling forward is better than stepping forward and then falling backward” — which is wrong for the actual goal of learning to walk.

This is a **credit assignment problem**: the total reward of a whole trajectory is being used to judge every single action inside it, even actions that happened early and had nothing to do with the final outcome.

Page 12 — Improving the Gradient (Reward-to-Go)

The key insight.

Policy behavior at time t does not affect rewards at time $t' < t$.

In plain English: **your action right now cannot change what already happened in the past.** If you fall at step 5, that cannot hurt or help your reward from steps 1–4. Those are already done. So it is unfair — and noisy — to judge an action at time t by the rewards that came **before** it. We should only hold an action responsible for the rewards that come **after** it (including the current moment). These future rewards are called the **reward-to-go**.

Old gradient (the flawed one).

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \right) \left(\sum_{t=1}^T r(s_t^i, a_t^i) \right)$$

The second sum is the **total reward of the entire trajectory** — it judges every action by the sum of all rewards, past and future.

New gradient (the fix).

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \left(\sum_{t'=t}^T r(s_{t'}^i, a_{t'}^i) \right)$$

Symbol	Meaning
$\sum_{t=1}^T$	Outer sum now goes over each individual timestep t , not just trajectories
$\nabla_{\theta} \log \pi_{\theta}(a_t^i s_t^i)$	Same as before — how to nudge θ for action a_t
$\sum_{t'=t}^T r(s_{t'}^i, a_{t'}^i)$	Reward-to-go: rewards from time t onward only

Instead of judging action a_t by the total reward of the whole episode, we only judge it by the rewards it could have actually influenced — everything from time t to the end. This is fairer and makes the gradient less noisy, because past rewards that had nothing to do with action a_t are no longer polluting its signal.

Back to the walking robot. With reward-to-go, τ^4 (step forward, then fall back) gets properly evaluated. The forward step at time $t = 1$ gets credited with the rewards from $t = 1$ onward — which include the positive forward velocity from that step. The backward fall at time $t = 2$ only gets blamed for rewards from $t = 2$ onward. The gradient now correctly separates the good action from the bad one within the same trajectory, instead of blurring them together under one total reward number.

1.11 Quiz: What the Gradient Does — The Walking Robot (Page 13)

The Setup

A humanoid robot is learning to walk. The reward $r(s, a)$ is the robot's forward speed. Move forward $\rightarrow$ positive reward. Move backward $\rightarrow$ negative reward.

Step 1 — Collect sample episodes

The robot tries 4 different episodes under the current policy $\pi_{\theta}(a_t | s_t)$:

Episode	What happened	Reward
τ^1	Falls forwards	Positive (moved forward)
τ^2	Slowly stumbles forwards	Positive
τ^3	Steadily walks forwards	Positive
τ^4	Runs forwards	Positive

All four got **positive reward** because all of them moved forward.

Step 2 — Compute the gradient

$$\nabla_{\theta} J(\theta) \approx \sum_i^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \left(\sum_{t'=t}^T r(s_{t'}^i, a_{t'}^i) \right)$$

- $\nabla_{\theta} J(\theta)$ — the gradient of the total reward with respect to the policy parameters; the “which direction to improve” signal
- $\sum_i^N$ — sum over all N sampled episodes

- $\sum_{t=1}^T$ — sum over every time step in an episode
- $\nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i)$ — how much each action’s probability should shift
- $\sum_{t'=t}^T r(s_{t'}^i, a_{t'}^i)$ — total future reward collected from time step t onwards (the “return”)

For every action taken in every episode, multiply “how much this action’s probability should change” by “how much reward followed from this point.” Average this over all episodes to get the gradient — the direction to nudge the policy.

The Surprising Answer

Since **all four episodes got positive reward**, the gradient gives a positive push to **all of them** — including the falling episode (τ^1). The gradient cannot tell the difference between “fall forward” and “run forward.” It just sees “positive reward $\rightarrow$ do more of that.” So the policy gets nudged toward falling forward, which is clearly not the ideal walking behavior.

This is the core problem: **policy gradient is high-variance** — its gradient estimates are noisy and sensitive to the overall scale of rewards. If everything gets positive reward, everything gets reinforced, even bad behavior.

1.12 Introducing the Baseline (Page 14)

The Core Idea

Instead of reinforcing an action whenever the reward is positive, reinforce it only when the reward is **better than average**. This is done by subtracting a **baseline** b from the reward.

The Modified Gradient Formula

$$\nabla_{\theta} J(\theta) = E_{\tau \sim p_{\theta}} [\nabla_{\theta} \log p_{\theta}(\tau) \cdot (r(\tau) - b)]$$

- $\nabla_{\theta} J(\theta)$ — gradient of the objective (direction to improve policy)
- $E_{\tau \sim p_{\theta}}[\cdot]$ — average over all episodes sampled from the current policy
- $\nabla_{\theta} \log p_{\theta}(\tau)$ — how much the probability of this episode’s actions should shift
- $r(\tau)$ — total reward of episode τ
- b — the baseline (a reference number we subtract)
- $(r(\tau) - b)$ — how much better or worse this episode was compared to the baseline

Instead of asking “did this episode get good reward?”, we now ask “did this episode do **better than the baseline**?” If yes, increase the probability of those actions. If no, decrease it. The baseline acts like a reference point — a “passing grade.”

What is the baseline?

The simplest choice is the **average reward** across all sampled episodes:

$$b = \frac{1}{N} \sum_i r(\tau)$$

- N — total number of sampled episodes
- $r(\tau)$ — reward of each episode
- b — their simple average

This is just computing the average score across all your robot’s attempts.

Applying it to the walking example

Say the average reward across the 4 episodes is some value b . Now:

- “Running forward” (τ^4) scores well above $b \rightarrow (r - b) > 0 \rightarrow$ get a **positive push** (do more of this)
- “Falling forward” (τ^1) scores below $b \rightarrow (r - b) < 0 \rightarrow$ get a **negative push** (do less of this)

Now the gradient correctly discourages falling and encourages running, even though both had positive raw rewards.

Does subtracting the baseline break anything?

No — and this is proved at the bottom of the slide. Subtracting a constant baseline leaves the **expected gradient unchanged** (it is unbiased). The proof shows:

$$E[\nabla_{\theta} \log p_{\theta}(\tau) \cdot b] = b \nabla_{\theta} \int p_{\theta}(\tau) d\tau = b \nabla_{\theta} \cdot 1 = 0$$

The key step: $\int p_{\theta}(\tau) d\tau = 1$ because all probabilities must sum to 1. The gradient of a constant (1) is zero. So the baseline contributes **zero** to the expected gradient — it does not bias the direction of learning at all. But it **reduces variance** significantly, because individual episode rewards are now measured relative to the average instead of in absolute terms. Smaller swings in the gradient = more stable learning.

1.13 Quiz: Sparse Rewards — Jacket Folding (Page 15)

The Setup

A robot is learning to fold a jacket. The reward has only three possible values:

Outcome	Reward
Doesn’t touch the jacket	0
Folded with wrinkles	0.5
Neatly folded	1

Step 1 — Sample episodes

Episode	What happened	Reward
τ^1	Doesn't touch the jacket	0
τ^2	Folds only the sleeves	≈ 0
τ^3	Flattens but doesn't fold	0
τ^4	Folds the jacket	1

Three out of four episodes got zero reward. Only τ^4 succeeded.

Step 2 — Compute the gradient with baseline

$$\nabla_{\theta} J(\theta) \approx \sum_i^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \left(\sum_{t'=t}^T r(s_{t'}^i, a_{t'}^i) - b \right)$$

- $\nabla_{\theta} J(\theta)$ — the gradient; tells us which direction to nudge the policy
- $\nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i)$ — how much each action's probability should shift
- $\sum_{t'=t}^T r(s_{t'}^i, a_{t'}^i)$ — total future reward from step t onward in episode i
- b — the baseline (average reward across all episodes, which is close to 0 here)

For each action taken in each episode, push that action's probability up or down based on how much better-than-average the outcome was. Subtract the baseline b so only above-average results get reinforced.

What happens in practice here?

Since most episodes scored 0 and the baseline $b \approx 0$, we get $(r - b) \approx 0$ for τ^1, τ^2, τ^3 — their gradient contributions are essentially **zero**. Only τ^4 , which scored 1, gives a meaningful signal.

So the gradient correctly points toward folding behavior — but the problem is it barely received any signal at all. Three out of four episodes contributed almost nothing to learning. This is called the **sparse reward problem**: when success is rare, most of your data is useless, the gradient is nearly flat most of the time, and learning is painfully slow and noisy. The baseline helped in the walking example (page 13) but did not fully fix this new problem.

1.14 Vanilla Policy Gradient and the On-Policy Problem (Page 16)

The Full Algorithm

1. **Sample** episodes $\{\tau^i\}$ from the current policy $\pi_{\theta}(a_t | s_t)$
2. **Compute** the gradient $\nabla_{\theta} J(\theta)$ using those episodes
3. **Update** the parameters
4. **Repeat** until the policy stops improving (convergence)

The update equation (Step 3):

$$\vec{\theta} = \vec{\theta} + \alpha \nabla_{\theta} J(\theta)$$

- $\vec{\theta}$ — the policy’s parameters (all the knobs controlling what actions the policy picks)
- α — the learning rate; how big a step to take (too big $\rightarrow$ overshoot, too small $\rightarrow$ crawl)
- $\nabla_{\theta} J(\theta)$ — the gradient; the direction that increases total reward

Take the current parameters $\vec{\theta}$, and nudge them in the direction of the gradient by a small amount α . Repeat. This is plain gradient ascent — climbing uphill on the reward surface one step at a time.

The Big Problem: On-Policy Learning

Vanilla policy gradient is **on-policy**, meaning the episodes used to compute the gradient in Step 2 must have been collected by the **exact current version of the policy**. Once you update θ in Step 3, those episodes are immediately outdated and thrown away. You can never reuse them.

This creates a painful cycle: collect data $\rightarrow$ compute gradient $\rightarrow$ update policy $\rightarrow$ **throw all data away** $\rightarrow$ collect new data $\rightarrow \dots$

Every single gradient step requires a fresh batch of episodes. If collecting data is slow (a real robot trying things in the physical world, or an expensive simulation), this wastes enormous amounts of time.

The natural question is: **what if we want to reuse data collected by an older version of the policy?** The answer is **off-policy RL using importance sampling** — a mathematical correction that lets you reuse old data by accounting for the fact that the old policy and the new policy chose actions differently. This is what the next slides address.

1.15 Off-Policy Policy Gradient (Page 17)

The Big Idea: Importance Sampling

Imagine you want to know what score a student would get on a different exam, but you only have their answers from the original exam. You can still estimate it — but you have to *adjust* each answer’s weight based on how different the two exams are. That adjustment factor is called the **importance sampling ratio**.

Same idea here: we collected episodes using the *old* policy θ , but we want to compute the gradient for a *new* (updated) policy θ' . We correct for this mismatch by multiplying by the ratio of the two policies’ probabilities.

Equation 1 — The Exact Off-Policy Gradient

$$\nabla_{\theta'} J(\theta') = \mathbb{E}_{\tau \sim p_{\theta}(\tau)} \left[\prod_{t=1}^T \frac{\pi_{\theta'}(a_t | s_t)}{\pi_{\theta}(a_t | s_t)} \left(\sum_{t=1}^T \nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t) \right) \left(\sum_{t'=t}^T r(s_{t'}, a_{t'}) - b \right) \right]$$

- θ' — the **new** policy’s parameters (the one we are trying to improve)
- θ — the **old** policy’s parameters (the one that actually collected the data)

- $\mathbb{E}_{\tau \sim p_{\theta}(\tau)}[\cdot]$ — average over episodes collected using the **old** policy θ
- $\prod_{t=1}^T \frac{\pi_{\theta'}(a_t | s_t)}{\pi_{\theta}(a_t | s_t)}$ — the **importance sampling correction**: a product of ratios, one per time step; each ratio asks “how much more or less likely would the new policy have been to take this same action?”
- $\nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t)$ — how much to shift each action’s probability under the new policy
- $\sum_{t'=t}^T r(s_{t'}, a_{t'}) - b$ — future reward from step t onward, minus the baseline

This is the REINFORCE gradient formula from before, but now with a correction factor bolted on — the big product in front. It uses data from the old policy, but multiplies each episode’s contribution by how differently the new policy would have behaved throughout that episode, step by step.

The Problem with Equation 1

The correction factor is a **product of T ratios** — one ratio multiplied together for every time step in the episode. If the episode is long (T is large), this product becomes unstable:

- If each ratio is slightly less than 1 $\rightarrow$ the product shrinks toward **zero** (gradient vanishes — no learning signal)
- If each ratio is slightly greater than 1 $\rightarrow$ the product explodes toward **infinity** (gradient blows up — wild, unstable updates)

This is like compounding interest but in a dangerous way. Multiplying together 100 numbers that are each 0.9 gives $0.9^{100} \approx 0.000027$ — basically nothing.

Equation 2 — Switching from Trajectories to Timesteps

$$\nabla_{\theta'} J(\theta') \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \frac{\pi_{\theta'}(s_{i,t}, a_{i,t})}{\pi_{\theta}(s_{i,t}, a_{i,t})} \nabla_{\theta'} \log \pi_{\theta'}(a_{i,t} | s_{i,t}) \left(\sum_{t'=t}^T r(s_{i,t'}, a_{i,t'}) - b \right)$$

- $\frac{\pi_{\theta'}(s_{i,t}, a_{i,t})}{\pi_{\theta}(s_{i,t}, a_{i,t})}$ — the ratio of how likely it was to **be in state s and take action a** under the new policy versus the old policy (a single ratio per timestep, not a product over all steps)
- $s_{i,t}, a_{i,t}$ — the state and action at time step t in episode i
- Everything else is the same as Equation 1

Instead of one giant product across the whole trajectory, we now use a **single ratio per time step** — comparing the new and old policy just at that one moment. This avoids the vanishing/exploding problem because there is no long chain of multiplications. The slide says this is “much less likely to explode/vanish.” However, the ratio $\frac{\pi_{\theta'}(s_{i,t}, a_{i,t})}{\pi_{\theta}(s_{i,t}, a_{i,t})}$ is the probability of both being in state s *and* taking action a . The state part is hard to compute — we do not know exactly how often the new policy would visit each state, because state visits depend on the full history of decisions. It is “hard to measure.”

Equation 3 — The Final Practical Form

$$\nabla_{\theta'} J(\theta') \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \frac{\pi_{\theta'}(a_{i,t} | s_{i,t})}{\pi_{\theta}(a_{i,t} | s_{i,t})} \nabla_{\theta'} \log \pi_{\theta'}(a_{i,t} | s_{i,t}) \left(\sum_{t'=t}^T r(s_{i,t'}, a_{i,t'}) - b \right)$$

- $\frac{\pi_{\theta'}(a_{i,t}|s_{i,t})}{\pi_{\theta}(a_{i,t}|s_{i,t})}$ — the ratio of **just the action probabilities**: given that we are already in state s , how much more or less likely is the new policy to pick action a compared to the old policy? The state part is dropped entirely.
- Everything else is the same as Equation 2

The state visitation ratio (the hard-to-measure part) is dropped, and only the action probability ratio is kept. This is an approximation — we are assuming the new and old policy visit states at roughly the same frequency — but it makes the formula practical. Both $\pi_{\theta'}(a | s)$ and $\pi_{\theta}(a | s)$ are just policy outputs that can be computed directly by plugging the state into the policy network.

Putting It All Together

The three equations represent a progression of trade-offs:

Equation	Correction factor	Explode/vanish?	Easy to compute?
Eq. 1	Product over all T steps	Yes, badly	Yes
Eq. 2	Single state-action ratio per step	No	No (state part is hard)
Eq. 3	Single action-only ratio per step	No	Yes

Equation 3 is the final usable formula. It lets you reuse data from an old policy, corrects for the mismatch using a simple per-step ratio, and is fully computable. This is the foundation for **off-policy** algorithms that are far more data-efficient than vanilla policy gradient.

1.16 Off-Policy Policy Gradient: Full Algorithm (Page 18)

The Algorithm (same three steps, different loop)

1. Sample episodes $\{\tau^i\}$ from policy $\pi_{\theta}(a_t | s_t)$
2. Compute $\nabla_{\theta} J(\theta)$ using those episodes
3. Update: $\vec{\theta} = \vec{\theta} + \alpha \nabla_{\theta} J(\theta)$

The update equation:

$$\vec{\theta} = \vec{\theta} + \alpha \nabla_{\theta} J(\theta)$$

- $\vec{\theta}$ — the policy's parameters (all the internal settings that control behavior)
- α — the learning rate; how big each step is
- $\nabla_{\theta} J(\theta)$ — the gradient; the direction that increases reward

Nudge the parameters in the direction that increases reward, scaled by the step size α . This is the same update as vanilla policy gradient.

What is actually different: the loop

The slide shows two arrows on the algorithm:

- **Red arrow** (on-policy loop): After step 3, loop back all the way to step 1 — collect **new** data before doing anything else.
- **Orange arrow** (off-policy loop): After step 3, loop back only to step 2 — recompute the gradient on the **same old data** using the importance sampling correction from page 17, and update again.

In vanilla policy gradient, after one update you must throw away all your data and go collect more. In off-policy policy gradient, you can reuse the same batch and do **multiple gradient steps** on it — because the importance sampling ratio in the gradient formula already corrects for the fact that the policy has changed since the data was collected.

Think of it like studying for an exam: on-policy is like solving one practice problem and then immediately going to get a new set — very inefficient. Off-policy is like solving the same set of practice problems multiple times, learning more from them each pass through.

1.17 The KL Divergence Constraint (Page 19)

The Problem

When you take many gradient steps on the same batch of old data, the new policy θ' can drift very far from the old policy θ that collected the data. But the importance sampling correction only works well when the two policies are *similar*. Once they diverge significantly:

- The old data no longer reflects the states the new policy would actually visit
- The gradient estimate becomes increasingly wrong
- You might make a bad update that seriously hurts the policy's performance — and you might not even realize it until it is too late

It is like navigating using an old map. A slightly outdated map is fine. But if the roads have completely changed, the map will lead you off a cliff.

The Fix — Constrain How Much the Policy Can Change

$$\mathbb{E}_{s \sim \pi_\theta} [D_{\text{KL}}(\pi_{\theta'}(\cdot | s) \parallel \pi_\theta(\cdot | s))] \leq \delta$$

- $\mathbb{E}_{s \sim \pi_\theta} [\cdot]$ — average over states that the **old** policy θ visited; we are measuring the difference in the situations the old policy actually experienced
- $D_{\text{KL}}(\cdot \parallel \cdot)$ — **KL divergence**: a mathematical ruler for measuring how different two probability distributions are; it equals 0 when they are identical, and grows larger as they diverge
- $\pi_{\theta'}(\cdot | s)$ — the **new** policy's action choices at state s (a full probability distribution over actions)
- $\pi_\theta(\cdot | s)$ — the **old** policy's action choices at state s
- δ — a small threshold (your “budget” for how much the policy is allowed to change); typically a small number like 0.01

On average across all the states the old policy visited, the new policy must not be “too different” from the old one. “Too different” is measured by KL divergence, and the constraint says it must stay at or below the budget δ . This limits how large each policy update can be, keeping the new policy close enough to the old one that the recycled data remains trustworthy.

Why This Matters

This constraint is the core idea behind **Trust Region Policy Optimization (TRPO)** — one of the most important algorithms in modern RL. The name comes from this: you trust the old data only within a certain region of policy space. Inside that region (where new and old policies are close), the gradient estimate is reliable. Outside it, you cannot trust the old data anymore and must stop.

A simpler version of this same idea later became **PPO (Proximal Policy Optimization)**, which enforces a similar limit without the full math of a trust region, and is today one of the most widely used RL algorithms in practice — including in training large language models like the ones behind ChatGPT.

2 Actor Critic Method

What is value function and what is Q function.

Value Function $V^\pi(s)$

What it answers: “How good is it to *be* in state s ?”

More precisely: if you’re currently in state s and you follow policy π from that point forward, what is the **expected total (sum of future) reward** you’ll collect?

$$V^\pi(s) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s \right]$$

It doesn’t care which action you take — it averages over whatever actions π would naturally choose. So $V^\pi(s)$ is purely a **state-level** signal.

Q-Function $Q^\pi(s, a)$

What it answers: “How good is it to be in state s , take a *specific* action a , and then follow π ?”

The key difference: you **commit to a specific action first**, then hand control back to the policy.

$$Q^\pi(s, a) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s, a_0 = a \right]$$

This is a **state-action level** signal — it lets you compare individual actions against each other.

The Key Relationship

$$V^\pi(s) = \mathbb{E}_{a \sim \pi(\cdot|s)} [Q^\pi(s, a)]$$

$V^\pi(s)$ is just $Q^\pi(s, a)$ **averaged over all actions weighted by the policy**. V sees the average picture; Q sees the picture for each individual action.

Why Both Exist

This distinction is what makes the **Advantage function** possible:

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

A^π tells you how much better (or worse) action a is compared to what the policy would do on average at state s . This is the core signal used by the **Critic** in Actor-Critic methods to give the Actor more informative feedback than raw rewards.

How V is Q averaged over all actions? Explain this with a concrete example.

Concrete Example

Say you're at state s , and your policy π has **3 possible actions** with these probabilities and Q -values:

Action	$\pi(a s)$ — probability	$Q^\pi(s, a)$ — expected future reward
a_1	0.5	10
a_2	0.3	4
a_3	0.2	1

What $V^\pi(s)$ Is

$V^\pi(s)$ asks: “If I just *run policy* π from state s — without picking any specific action myself — what reward do I expect?”

Since π will pick a_1 50% of the time, a_2 30%, a_3 20%, the expected reward is just the **probability-weighted sum of Q -values**:

$$V^\pi(s) = 0.5 \times 10 + 0.3 \times 4 + 0.2 \times 1 = 5 + 1.2 + 0.2 = \mathbf{6.4}$$

That's it. V doesn't have any secret information — it's literally just averaging $Q^\pi(s, a)$ weighted by how often the policy *would* choose each action.

Why This Makes Intuitive Sense

- $Q^\pi(s, a_1) = 10 \rightarrow$ if you *commit* to a_1 , you expect 10.
- But you don't always take a_1 . Half the time you do, and the other half you take worse actions.
- So the **actual expected reward from state s** — without fixing an action — is 6.4, which is lower than the best possible Q .

$V^\pi(s) = 6.4$ is the **average performance of the policy at state s** . $Q^\pi(s, a)$ is the performance of **one specific choice**.

The Advantage Connection

This is why the Advantage function is meaningful:

- $A^\pi(s, a_1) = 10 - 6.4 = +3.6 \rightarrow a_1$ is better than average
- $A^\pi(s, a_2) = 4 - 6.4 = -2.4 \rightarrow a_2$ is worse than average
- $A^\pi(s, a_3) = 1 - 6.4 = -5.4 \rightarrow a_3$ is much worse than average

V acts as the **baseline**. Q minus that baseline tells you which actions are actually worth doing more often — which is exactly the signal the Actor uses to improve.

What does A tell us?

$A^\pi(s, a)$ tells you: “**Is this action better or worse than what the policy would do on average at this state — and by how much?**”

- $A > 0 \rightarrow$ this action is better than average. Do it more.
- $A < 0 \rightarrow$ this action is worse than average. Do it less.
- $A = 0 \rightarrow$ this action is exactly as good as following the policy blindly.

From the previous example:

- a_1 has $A = +3.6 \rightarrow$ push the policy to favor a_1 more
- a_2 has $A = -2.4 \rightarrow$ push the policy away from a_2
- a_3 has $A = -5.4 \rightarrow$ strongly push away from a_3

The key insight: **raw reward tells you if something was good in absolute terms. Advantage tells you if it was good *relative to your other options***. That relative comparison is what makes the gradient update meaningful — instead of just chasing high-reward actions, the agent learns which actions are genuinely better than its current behavior.

Issues with Policy Gradient Methods

The Big Idea

This slide is asking: “**What’s wrong with how the AI currently learns?**” The answer is — it’s clumsy. It blames or rewards the *entire sequence of actions* together, even when only one action was actually good or bad. This causes two big problems, shown through two examples.

The Equation

$$\nabla_\theta J(\theta) \approx \sum_i \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t^i | s_t^i) \left(\sum_{t'=t}^T r(s_{t'}, a_{t'}^i) - b \right)$$

What each symbol means:

Symbol	What it is
$\nabla_{\theta} J(\theta)$	The direction to tweak the policy to get better rewards
θ	The knobs inside the AI's brain (its neural network weights)
$\sum_i$	"Add this up across all tries (trajectories)"
$\sum_t$	"Add this up across every time step in one try"
$\log \pi_{\theta}(a_t^i s_t^i)$	How likely the policy thought action a was in state s
$\sum_{t'=t} r(\dots)$	Total reward collected from this moment forward (reward to go)
b	A baseline — the average reward, subtracted to reduce noise

What the equation is saying, in plain English:

For every action the AI took, look at all the rewards it collected *after* that action. If those future rewards were high, nudge the AI to make that action more likely next time. If they were low, nudge it to avoid that action. Subtract b so you're comparing to the average, not reacting to raw numbers.

Problem 1: The Walking Robot

The robot takes **one small step forward**, then **falls backwards**. The total reward for the trip is negative (it fell!).

The equation sees the negative total reward and says: "*That step forward must have been bad — do it less!*"

But that's **wrong**. The step forward was actually a good move. The fall was the problem. The algorithm can't tell the difference — it punishes the whole sequence including the good part.

- × It pushes down the chance of "step forward" even though stepping forward was fine.
- × It wastes the data because it learned the wrong lesson.

Problem 2: The Jacket-Folding Robot

Three different practice runs (trajectories):

- τ^2 — folds only the sleeves
- τ^3 — flattens the jacket only
- τ^4 — fully folds the jacket (gets reward)

Only τ^4 gets a reward of 1. τ^2 and τ^3 get **zero reward** because the task wasn't fully done.

The algorithm sees zero reward for τ^2 and τ^3 and says: "*Nothing useful happened there, ignore it.*"

But that's **also wrong**. Folding the sleeves and flattening the jacket are *progress* — they were good intermediate steps that the AI should learn to repeat.

- × With sparse rewards, the AI wastes most of its practice runs.
- ✓ The question the slide ends on: **Can we teach the AI to know what's good and bad, even mid-task?** — That's exactly what Actor-Critic methods solve.

Improving Policy Gradients (Page 5)

The starting equation (from last time):

$$\nabla_{\theta} J(\theta) \approx \sum_i \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \left(\sum_{t'=t}^T r(s_{t'}^i, a_{t'}^i) - b \right)$$

Symbol	What it means
$\nabla_{\theta} J(\theta)$	The direction to improve the policy — like a compass saying “adjust your brain th
θ	The AI’s brain settings (numbers inside the neural network)
$\sum_i$	Add this up across all practice runs
$\sum_{t=1}^T$	Add this up across every step of one run
$\nabla_{\theta} \log \pi_{\theta}(a_t^i s_t^i)$	How confident the AI was in the action it chose at that step
$\sum_{t'=t}^T r(\dots)$	All the rewards collected <i>from this step forward</i> (called reward to go)
b	A baseline — the average reward — subtracted to reduce noise

The “reward to go” part is trying to answer: “*If I took this action in this situation, how much total reward did I collect afterward?*” That quantity — the expected future reward of taking action a from state s — is exactly the definition of the **Q-function**. So the slide says:

Step 1 — Reward to go equals the Q-function:

$$\sum_{t=t'}^T \mathbb{E}_{\pi_{\theta}}[r(s_{t'}, a_{t'} | s_t, a_t)] = Q_{\pi}(s_t, a_t)$$

Symbol	What it means
$\mathbb{E}_{\pi_{\theta}}[\cdot]$	The average outcome if you repeat this many times under policy π
$r(s_{t'}, a_{t'})$	Reward you get at each future step
$Q_{\pi}(s_t, a_t)$	The Q-function: expected total future reward for taking action a_t in state s_t

Instead of using the messy, noisy actual rewards from one run, we can use the Q-function — which is the *average* expected reward. The diagram on the slide shows this visually: from state s , you take action a , then let policy π run freely. All those possible future paths spread out like branches, and Q is the expected sum of rewards across all of them. One single run (the “reward to go”) is just one noisy sample of Q.

Step 2 — The baseline b turns out to be the Value function:

$$b = \text{average reward} = \frac{1}{N} \sum_i Q_{\pi}(s_{i,t}, a_{i,t}) = V_{\pi}(s_t)$$

Symbol	What it means
$\frac{1}{N} \sum_i$	Average across all N practice runs
$Q_{\pi}(s_{i,t}, a_{i,t})$	Q-value for each run’s action at this step
$V_{\pi}(s_t)$	The Value function — the average Q across all possible actions the policy would take

The baseline b was supposed to be “the average reward” to compare against. When you average Q over all the actions the policy would naturally take, you get exactly $V_\pi(s_t)$ — the Value function. So b isn’t arbitrary; it’s the average quality of the current state.

Step 3 — The whole thing simplifies to the Advantage:

$$\nabla_\theta J(\theta) \approx \sum_i \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t^i | s_t^i) A_\pi(s_{i,t}, a_{i,t})$$

Symbol	What it means
$A_\pi(s_{i,t}, a_{i,t})$	The Advantage — how much better this specific action was compared to the average

Since $Q - b = Q - V = A$ (the Advantage function), the whole messy reward-to-go-minus-baseline expression collapses into just A_π . The slide is saying: instead of asking “did I get good rewards?”, ask “was this action *better than average* for this situation?” — that’s a much smarter, cleaner signal for improving the policy. Actions better than average get a positive score $\rightarrow$ do them more. Actions worse than average get a negative score $\rightarrow$ do them less.

How using A is different from using `reward_to_go - b`? `reward_to_go = Q` and `b = V`. I see no difference.

You’re right — mathematically they are **identical**. $A_\pi = Q - V = \text{reward_to_go} - b$, by definition. No difference on paper.

The difference is in **how you compute them in practice**.

`reward_to_go - b` (pure policy gradient way):

You run a full trajectory, collect actual rewards, sum them up. That sum *is* your estimate of Q . You subtract the average across runs as b . Both numbers come from **raw sampled rewards** — noisy, because any single run is just one random path through the environment.

A via the Critic (actor-critic way):

You train a **neural network** to learn $V(s)$. Then you compute the advantage as:

$$\hat{A}_t = r_t + V_\phi(s_{t+1}) - V_\phi(s_t)$$

You’re no longer summing up a long noisy chain of rewards. The network gives you a smooth, learned estimate of value at any state.

The practical consequence:

	<code>reward_to_go - b</code>	Critic-estimated A
Source	Raw rewards from full runs	Neural network estimates
Variance	High	Lower
Bias	Unbiased	Can be biased if network is wrong
Data needed	Many full trajectories	Can learn from shorter, incomplete ones

So the slide isn't introducing a new *math* — it's showing that A is the right **target concept**, and the whole point of actor-critic is to **estimate A better** than raw reward sums allow. The Critic exists precisely to give you a lower-variance estimate of A than you'd get by just running trajectories.

Online Reinforcement Learning (Page 6)

This slide shows the **big picture loop** of how the AI actually trains. It has three stages that repeat forever.

Step 0 — Start somewhere:

The policy π is initialized — either randomly (the AI does dumb stuff at first), by copying a human expert (imitation learning), or using some hand-crafted rules (heuristics).

Then the loop runs:

Stage 1 → Run policy to collect data.

Let the AI play in the environment for a while. Record what happened: which states it visited, which actions it took, what rewards it got. This is the raw experience batch.

Stage 2 → Fit a model to estimate expected return.

Use that collected data to train a neural network to estimate V^π , Q^π , or A^π — i.e., teach the Critic how good each state or action was. This is the “how good was that?” step.

Stage 3 → Improve the policy:

$$\theta \leftarrow \theta + \nabla_\theta J(\theta)$$

Symbol	What it means
θ	The AI's brain settings (neural network weights)
$\leftarrow$	“Update this to become”
$\nabla_\theta J(\theta)$	The direction that makes the policy better

Take a step in the direction that improves the policy, using the Advantage estimates from Stage 2. Then go back to Stage 1 with the new, slightly better policy. Repeat.

Estimating Expected Return (Page 7)

This slide answers: how do we actually compute A_π ? We start from its definition and simplify it into something we can calculate.

Step 1 — Definition of Advantage:

$$A_\pi(s_t, a_t) = Q_\pi(s_t, a_t) - V_\pi(s_t)$$

Symbol	What it means
$A_\pi(s_t, a_t)$	How much better action a_t is compared to the average action at state s_t
$Q_\pi(s_t, a_t)$	Expected total future reward if you take action a_t at state s_t , then follow π
$V_\pi(s_t)$	Expected total future reward at state s_t if you just follow π normally

Advantage is simply Q minus V — how much better (or worse) your specific action was compared to what the policy would do on average.

Step 2 — Expand Q:

$$Q_\pi(s_t, a_t) = \sum_{t'=t}^T \mathbb{E}_{\pi_\theta}[r(s_{t'}, a_{t'} \mid s_t, a_t)]$$

This is just the definition of Q — the expected sum of all future rewards from this step onward. Nothing new here, just spelling it out.

Step 3 — Rewrite Q using just one step + V:

$$Q_\pi(s_t, a_t) = r(s_t, a_t) + \mathbb{E}_{s_{t+1} \sim p(\cdot \mid s_t, a_t)}[V_\pi(s_{t+1})]$$

Symbol	What it means
$r(s_t, a_t)$	The immediate reward you get right now
$\mathbb{E}_{s_{t+1} \sim p(\cdot \mid s_t, a_t)}[\cdot]$	The average over all possible next states you could land in
$V_\pi(s_{t+1})$	The value of wherever you end up next

Instead of summing rewards all the way to the end of the episode, we can split it: the reward you get *right now*, plus the value of the *next state* (which already accounts for everything after). This is a much shorter calculation.

Step 4 — Approximate by using the actual next state:

$$Q_\pi(s_t, a_t) \approx r(s_t, a_t) + V_\pi(s_{t+1})$$

We drop the expectation and just use the real s_{t+1} that actually happened in the run. This is an approximation — one sample instead of averaging over all possible next states — but it works well enough in practice.

Step 5 — Final practical formula for Advantage:

$$A_\pi(s_t, a_t) \approx r(s_t, a_t) + V_\pi(s_{t+1}) - V_\pi(s_t)$$

Symbol	What it means
$r(s_t, a_t)$	Reward you got right now
$V_\pi(s_{t+1})$	How good the next state is (estimated by the Critic network)
$V_\pi(s_t)$	How good the current state was expected to be (estimated by the Critic network)

This is the key result. To compute A , you only need: the reward you just got, plus what the Critic says about the next state, minus what the Critic says about the current state. No long reward chains. No waiting until the end of the episode. This difference $r + V(s_{t+1}) - V(s_t)$ is called the **TD error** (Temporal Difference error). The slide closes by asking: but how do we estimate $V_\pi(s_t)$? — that's answered on the next slides using Monte Carlo or TD learning.

Actor-Critic Algorithm (Page 8)

This slide puts it all together as a concrete step-by-step algorithm. Two networks are running: the **Actor** (the policy, picks actions) and the **Critic** (the value function, judges states).

Step 1 — Sample a batch of data:

$$\{(s_{1,i}, a_{1,i}, \dots, s_{T,i}, a_{T,i})\} \text{ from } \pi_\theta$$

Run the current policy in the environment and record everything that happened across N trajectories: every state visited, every action taken, every reward received.

Step 2 — Fit the Critic:

Train the value network $\hat{V}_\phi^{\pi_\theta}$ on the collected data — i.e., teach it to predict the total future reward from any state it saw. ϕ are the Critic's own internal weights, separate from the Actor's θ .

Step 3 — Compute the Advantage estimate:

$$\hat{A}^{\pi_\theta}(s_{i,t}, a_{i,t}) = r(s_{i,t}, a_{i,t}) + \hat{V}_\phi^{\pi_\theta}(s_{i,t+1}) - \hat{V}_\phi^{\pi_\theta}(s_{i,t})$$

Symbol	What it means
$\hat{A}^{\pi_\theta}$	Estimated advantage — the hat $\hat$ means it's an approximation, not exact
$r(s_{i,t}, a_{i,t})$	Actual reward received at this step
$\hat{V}_\phi^{\pi_\theta}(s_{i,t+1})$	Critic's estimate of how good the <i>next</i> state is
$\hat{V}_\phi^{\pi_\theta}(s_{i,t})$	Critic's estimate of how good the <i>current</i> state is

This is the TD error from page 7 — immediate reward plus next-state value minus current-state value. Positive means this action was better than expected; negative means it was worse.

Step 4 — Compute the policy gradient:

$$\nabla_\theta J(\theta) \approx \sum_i^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t^i | s_t^i) \hat{A}^{\pi_\theta}(s_{i,t}, a_{i,t})$$

Same formula as before — but now the Advantage is estimated by the Critic network, not from raw noisy reward sums.

Step 5 — Update the Actor:

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$$

Symbol	What it means
α	Learning rate — how big a step to take (a small number like 0.001)
$\nabla_{\theta} J(\theta)$	The direction that improves the policy

Nudge the Actor’s weights in the direction that makes good-advantage actions more likely. Then go back to Step 1 with the updated policy.

The diagram on the right shows the data flow: the environment sends a *state* to both the Actor and Critic. The Actor outputs an *action*. The environment responds with a *reward* and a new state. The Critic uses the reward and both states to compute the **TD error**, which is used to update itself (Critic training) and to provide the Advantage signal back to the Actor (Actor training).

Off-Policy Actor-Critic: Version 1 (Multiple Gradient Steps) & Too Many Gradient Steps

The Big Idea

Imagine a robot learning to play a video game. It tries some moves, records what happened, and then uses that recording to get smarter. But here’s the trick — it tries to learn *a lot* from that one recording before throwing it away. That’s what this page is about.

The algorithm has 5 steps:

Step 1 — Collect experiences. The robot (called the “actor,” running policy π_{θ}) plays through the game several times and records all the situations (states) and moves (actions) it made.

Step 2 — Train the Critic. A second brain called the “critic” ($\hat{V}_{\phi}$) looks at all those recordings and learns to estimate: “*From this situation, how much total reward can we expect?*”

Step 3 — Calculate the Advantage.

$$\hat{A}^{\pi_{\theta}}(s_{i,t}, a_{i,t}) = r(s_{i,t}, a_{i,t}) + \hat{V}_{\phi}^{\pi_{\theta}}(s_{i,t+1}) - \hat{V}_{\phi}^{\pi_{\theta}}(s_{i,t})$$

What each symbol means:

- $\hat{A}^{\pi_{\theta}}$ — The “advantage.” How much *better* (or worse) was this action compared to what we expected?
- $s_{i,t}$ — The situation (state) the robot was in, at time step t , in episode i .
- $a_{i,t}$ — The action the robot took in that situation.

- $r(s_{i,t}, a_{i,t})$ — The immediate reward earned right after taking that action.
- $\hat{V}_\phi^{\pi_\theta}(s_{i,t+1})$ — The critic’s estimate of future rewards from the *next* situation.
- $\hat{V}_\phi^{\pi_\theta}(s_{i,t})$ — The critic’s estimate of future rewards from the *current* situation.

What it’s saying: The advantage is like asking: “*Was this move smarter than average?*” You take the reward you actually got, plus how good the future looks from where you landed, then subtract how good the future was already expected to look from where you started. If that number is positive, the move was better than expected — a smart move. If negative, it was worse. Think of it like getting a surprise quiz score — the advantage is how much you beat (or missed) your expected grade.

Step 4 — Calculate how to improve the policy.

$$\nabla_{\theta'} J(\theta') \approx \sum_i \sum_{t=1}^T \frac{\pi_{\theta'}(a_{i,t}|s_{i,t})}{\pi_\theta(a_{i,t}|s_{i,t})} \nabla_{\theta'} \log \pi_{\theta'}(a_t^i | s_t^i) \hat{A}^{\pi_\theta}(s_{i,t}, a_{i,t})$$

What each symbol means:

- $\nabla_{\theta'} J(\theta')$ — The direction to nudge the new policy θ' to make it better (the gradient).
- $\pi_{\theta'}(a_{i,t}|s_{i,t})$ — The probability the *new* policy would pick that same action in that situation.
- $\pi_\theta(a_{i,t}|s_{i,t})$ — The probability the *old* policy picked that action (the one that collected the data).
- The fraction — The *importance weight*. It’s a correction factor.
- $\nabla_{\theta'} \log \pi_{\theta'}(\cdot)$ — How sensitive the new policy is to small changes, for that particular action.
- $\hat{A}^{\pi_\theta}(\cdot)$ — The advantage we computed in Step 3.
- $\sum_i \sum_{t=1}^T$ — Sum this up across all episodes (i) and all time steps ($t = 1$ to T).

What it’s saying: We collected data using the *old* policy, but we want to improve the *new* policy. The problem is the data was collected under different rules than the ones we’re now updating. So we use the fraction (importance weight) as a *correction factor* — like converting currencies. If the new policy would have chosen that action much more often than the old policy did, the weight is high, meaning that action gets more credit. The whole equation then pushes the new policy to do more of the things that had a high advantage, adjusted fairly for the difference between old and new behavior.

Step 5 — Update the policy. Move the policy in the direction that makes it better, scaled by a learning rate α :

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$$

The Warning: If you do Steps 4–5 too many times on the same old data, things break.

The Problem — Too Many Gradient Steps

Imagine you studied for a test using last week’s notes, but the test is actually on new material. The more you study those old notes, the more confident you get — but in the *wrong* direction. The same thing happens here: if you update the policy too many times using old data, the new policy drifts so far from the old one that the importance weights become useless lies. The new policy “overfits” to the old data.

Fix 1 — Put a Leash on the Policy (KL Constraint)

$$\mathbb{E}_{s \sim \pi_\theta} [D_{\text{KL}}(\pi_{\theta'}(\cdot | s) \parallel \pi_\theta(\cdot | s))] \leq \delta$$

What each symbol means:

- $\mathbb{E}_{s \sim \pi_\theta}$ — The average, taken over all situations s that the old policy π_θ would visit.
- $D_{\text{KL}}(\cdot \parallel \cdot)$ — KL divergence: a way to measure *how different* two probability distributions are. Think of it as the “distance” between two sets of rules.
- $\pi_{\theta'}(\cdot | s)$ — The new policy’s behavior in situation s .
- $\pi_\theta(\cdot | s)$ — The old policy’s behavior in situation s .
- $\leq \delta$ — The difference must stay below a small limit δ (delta), a number you set in advance.

What it’s saying: This equation is a rule that says: “*The new policy is not allowed to become too different from the old policy.*” You measure how different they are using KL divergence — a mathematical way of comparing two sets of decisions. If the new policy starts changing too drastically, you stop it. It’s like telling a student: “*You can revise your answers, but don’t change more than 10% of them.*” This idea is used in training modern AI language models (like ChatGPT) to make sure they don’t go off the rails during fine-tuning.

Fix 2 — Clip the Importance Weights (The PPO Idea)

Instead of directly measuring how different the policies are, what if we just *cap* the importance weight ($\pi_{\theta'}/\pi_\theta$) so it can never get too large or too small? This doesn’t stop the policy from changing, but it removes the *reward* for changing too much — the policy has no reason to drift far because it can’t exploit extreme importance weights anymore. This is the core idea behind **PPO (Proximal Policy Optimization)**, one of the most popular and widely used training algorithms in modern AI, including the training of large language models.

Surrogate Objective

The Big Idea

This page answers one smart question: “*Instead of computing a complicated gradient every time, can we find a simpler score to just maximize directly?*” The answer is yes — and that simpler score is called the **surrogate objective**.

Step 1 — Start with the gradient from page 9

$$\nabla_{\theta'} J(\theta') \approx \sum_i \sum_{t=1}^T \frac{\pi_{\theta'}(a_{i,t} | s_{i,t})}{\pi_{\theta}(a_{i,t} | s_{i,t})} \nabla_{\theta'} \log \pi_{\theta'}(a_t^i | s_t^i) \hat{A}^{\pi_{\theta}}(s_{i,t}, a_{i,t})$$

This was explained on page 9. The two key pieces being looked at more carefully here are: the importance weight fraction, and the log-gradient of the policy.

Step 2 — A Useful Math Trick

$$p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau) = \nabla_{\theta} p_{\theta}(\tau)$$

What each symbol means:

- $p_{\theta}(\tau)$ — The probability of a sequence of events τ (like a full episode of play) under policy with parameters θ .
- $\nabla_{\theta} \log p_{\theta}(\tau)$ — How sensitive the *log* of that probability is to small changes in θ .
- $\nabla_{\theta} p_{\theta}(\tau)$ — How sensitive the *actual* probability is to small changes in θ .

What it's saying: This is a well-known math shortcut called the **log-derivative trick**. It says: multiplying a probability by the gradient of its own log is the same as just taking the gradient of the probability directly. Think of it like this: if you know the gradient of the “log score” and you also know the score itself, you can recover the gradient of the score without doing extra work. This lets us collapse the importance weight and the log-gradient into a single, cleaner expression — the gradient of the probability ratio. It's a pure algebra trick that rewrites one messy form into another simpler form.

Step 3 — The Surrogate Objective

Applying that trick to the gradient equation allows us to “un-wrap” the log and arrive at a plain function — not a gradient, but a score you can just maximize directly:

$$\tilde{J}(\theta') \approx \sum_{t,i} \frac{\pi_{\theta'}(a_{i,t} | s_{i,t})}{\pi_{\theta}(a_{i,t} | s_{i,t})} \hat{A}^{\pi_{\theta}}(s_{t,i}, a_{t,i})$$

What each symbol means:

- $\tilde{J}(\theta')$ — The surrogate objective. A fake (but useful) score that stands in for the real objective $J(\theta')$. The tilde signals it's a substitute.
- $\pi_{\theta'}(a_{i,t} | s_{i,t})$ — The probability that the *new* policy picks action a in situation s .
- $\pi_{\theta}(a_{i,t} | s_{i,t})$ — The probability that the *old* policy picked that same action.
- The fraction — The importance weight, correcting for the fact that data came from the old policy.
- $\hat{A}^{\pi_{\theta}}(s_{t,i}, a_{t,i})$ — The advantage: how much better that action was compared to average.

- $\sum_{t,i}$ — Sum over all time steps t and all episodes i .

What it’s saying: The surrogate objective $\tilde{J}(\theta')$ is a simplified, easy-to-maximize score. Instead of computing a complex gradient formula at every step, you just directly maximize this expression. It asks: *“For every action taken in the old recordings, how much more (or less) likely would the new policy be to make that same move, and was that move a good one?”* If the new policy gives higher probability to good moves (high advantage) and lower probability to bad ones (negative advantage), the score goes up. The key insight is that when the new policy is still close to the old one ($\theta' \approx \theta$), taking the gradient of $\tilde{J}$ gives you the exact same result as taking the gradient of the real objective J . So you can optimize $\tilde{J}$ as a stand-in — it’s mathematically equivalent at the point that matters, and much easier to work with.

The Two Warnings at the Bottom

- **Good news:** Maximizing $\tilde{J}$ naturally pushes the policy to take good actions more often — exactly what we want.
- **Bad news:** Because $\tilde{J}$ rewards giving higher probability to good actions, the policy is tempted to change *a lot* — drifting far from the old policy. If you run too many gradient steps on $\tilde{J}$ with the same old data, the ratio $\pi_{\theta'}/\pi_{\theta}$ grows huge and meaningless, and the policy overfits. This is precisely the “too many gradient steps” problem from page 10, and why PPO and KL constraints exist.

How did the log part of the initial equation vanish?

The log vanishes because of the math trick shown in the middle of the page. Here’s exactly how, step by step.

Look at the two boxed parts in Equation 1 together:

$$\frac{\pi_{\theta'}(a|s)}{\pi_{\theta}(a|s)} \cdot \nabla_{\theta'} \log \pi_{\theta'}(a|s)$$

The math trick says:

$$p(\tau) \cdot \nabla \log p(\tau) = \nabla p(\tau)$$

Apply that to the numerator only — $\pi_{\theta'}(a|s)$ is sitting in the fraction AND inside the log. So:

$$\pi_{\theta'}(a|s) \cdot \nabla_{\theta'} \log \pi_{\theta'}(a|s) = \nabla_{\theta'} \pi_{\theta'}(a|s)$$

Plug that back in:

$$\frac{\pi_{\theta'}(a|s) \cdot \nabla_{\theta'} \log \pi_{\theta'}(a|s)}{\pi_{\theta}(a|s)} = \frac{\nabla_{\theta'} \pi_{\theta'}(a|s)}{\pi_{\theta}(a|s)}$$

Since π_{θ} (the old policy) **doesn’t depend on** θ' , it’s just a constant with respect to the gradient. So you can pull $\nabla_{\theta'}$ outside:

$$= \nabla_{\theta'} \left[\frac{\pi_{\theta'}(a|s)}{\pi_{\theta}(a|s)} \right]$$

Now the full equation becomes:

$$\nabla_{\theta'} J(\theta') = \nabla_{\theta'} \sum_{t,i} \frac{\pi_{\theta'}(a_{t,i} | s_{t,i})}{\pi_{\theta}(a_{t,i} | s_{t,i})} \hat{A}^{\pi_{\theta}}(s_{t,i}, a_{t,i})$$

The right-hand side is just $\nabla_{\theta'}$ acting on $\tilde{J}(\theta')$. So:

$$\nabla_{\theta'} J(\theta') = \nabla_{\theta'} \tilde{J}(\theta')$$

The log didn't “vanish” — it got absorbed. The log-gradient and the importance weight were multiplied together, and the trick collapsed that product into a single clean gradient of the ratio. What's left inside the gradient sign is exactly $\tilde{J}$. Since the gradients match, you can just maximize $\tilde{J}$ directly instead of dealing with the log form.

Pages 12, 13, 14 — Proximal Policy Optimization (PPO)

All three pages are about one algorithm: PPO. Think of PPO as a smarter, safer way to train an AI agent.

Page 12 — PPO's First Two Tricks

Before the tricks, here is the starting point (the surrogate objective):

$$\tilde{J}(\theta') \approx \sum_{t,i} \frac{\pi_{\theta'}(a_{t,i} | s_{t,i})}{\pi_{\theta}(a_{t,i} | s_{t,i})} \hat{A}^{\pi_{\theta}}(s_{t,i}, a_{t,i})$$

What each symbol means:

- $\tilde{J}(\theta')$ — the score we are trying to maximize for the new policy
- θ' — the parameters (settings) of the new policy we are training
- θ — the parameters of the old policy we collected data with
- $\pi_{\theta'}(a_{t,i} | s_{t,i})$ — probability the new policy picks action a in situation s
- $\pi_{\theta}(a_{t,i} | s_{t,i})$ — probability the old policy picks the same action
- $\hat{A}^{\pi_{\theta}}(s_{t,i}, a_{t,i})$ — the advantage: how much better this action was compared to average
- $\sum_{t,i}$ — add this up across all time steps t and all collected episodes i

What this equation is saying: We collected experience using the old policy θ , but we want to improve our new policy θ' without collecting new data each time. The fraction $\frac{\pi_{\theta'}}{\pi_{\theta}}$ is called the importance weight — it corrects for the fact that data was collected under a different policy. If the new policy would have chosen the same action even more often, this ratio is big; if less often, it is small. We multiply this ratio by the advantage $\hat{A}$, which tells us whether that action was actually good or bad. If an action was good (positive advantage) AND the new policy would do it more (ratio > 1), the score goes up. If the new policy has drifted far from the old one, the ratio explodes to huge or tiny numbers — that is the problem this equation alone cannot fix, which is why we need the tricks.

Trick #1 — Clip the importance weights

$$\tilde{J}(\theta') \approx \sum_{t,i} \text{clip}\left(\frac{\pi_{\theta'}(a_{t,i} \mid s_{t,i})}{\pi_{\theta}(a_{t,i} \mid s_{t,i})}, 1 - \epsilon, 1 + \epsilon\right) \hat{A}^{\pi_{\theta}}(s_{t,i}, a_{t,i})$$

New symbols:

- $\text{clip}(\cdot, 1 - \epsilon, 1 + \epsilon)$ — a safety clamp: if the value goes below $1 - \epsilon$, force it to $1 - \epsilon$; if it exceeds $1 + \epsilon$, force it to $1 + \epsilon$
- ϵ (epsilon) — a small number (typically 0.1 or 0.2) that sets how far the ratio is allowed to stray from 1

What this equation is saying: The importance weight ratio can get dangerously large or small when the new and old policies differ a lot, making training unstable. Clipping simply forces that ratio to stay inside the window $[1 - \epsilon, 1 + \epsilon]$ — for example $[0.8, 1.2]$ if $\epsilon = 0.2$. If the new policy has already changed too much from the old one, we stop rewarding it for changing further. It is like saying “you have wandered far enough from where you started — no extra credit for going further.” This keeps training stable.

Trick #2 — Take the minimum

$$\tilde{J}(\theta') \approx \sum_{t,i} \min\left(\frac{\pi_{\theta'}}{\pi_{\theta}} \hat{A}, \text{clip}\left(\frac{\pi_{\theta'}}{\pi_{\theta}}, 1 - \epsilon, 1 + \epsilon\right) \hat{A}\right)$$

New symbols:

- $\min(\cdot, \cdot)$ — take whichever of the two values is smaller

What this equation is saying: Trick #1 clipped the ratio, but that alone has a loophole: when the advantage is negative (meaning the action was bad), clipping might accidentally still give the policy an incentive to make changes. Taking the minimum of the clipped and unclipped objective plugs that loophole. Imagine two numbers: the “raw” score and the “clamped” score. We always take the worse one. This makes PPO pessimistic by design — it never rewards the policy for making a change that is not clearly safe and beneficial. This is the core PPO objective.

Page 13 — Trick #3: Generalized Advantage Estimation (GAE)

N-step Advantage

$$\hat{A}_n^{\pi}(s_t, a_t) = \sum_{t'=t}^{t+n} \gamma^{t'-t} r(s_{t'}, a_{t'}) - \hat{V}_{\phi}^{\pi}(s_t) + \gamma^n \hat{V}_{\phi}^{\pi}(s_{t+n})$$

What each symbol means:

- $\hat{A}_n^{\pi}(s_t, a_t)$ — the estimated advantage of taking action a_t at state s_t , looking n steps ahead
- $\sum_{t'=t}^{t+n}$ — sum rewards from now (t) up to n steps in the future

- $\gamma^{t'-t}$ — discount factor: rewards further in the future count less (closer = worth more)
- $r(s_{t'}, a_{t'})$ — actual reward received at each step
- $\hat{V}_\phi^\pi(s_t)$ — the critic’s estimated value of the current state (what it expected)
- $\gamma^n \hat{V}_\phi^\pi(s_{t+n})$ — the discounted value estimate at the state n steps later (a bootstrapped guess of what comes after)

What this equation is saying: Instead of looking either one step ahead (fast but inaccurate) or all the way to the end (accurate but noisy), this looks exactly n steps ahead. You sum up the real rewards you collected for n steps, then add the critic’s guess of what happens after step n (instead of running all the way to the end). Subtracting $\hat{V}(s_t)$ turns this from a raw return into an advantage — how much better or worse things went compared to what the critic expected. Small n = lower variance but possibly biased; large n = more accurate but noisier. GAE blends all values of n together.

GAE (Generalized Advantage Estimation)

$$\hat{A}_{\text{GAE}}^\pi(s_t, a_t) = \sum_{n=1}^{\infty} w_n \hat{A}_n^\pi(s_t, a_t), \quad w_n \propto \lambda^{n-1}$$

What each symbol means:

- $\hat{A}_{\text{GAE}}^\pi$ — the final blended advantage estimate
- $\sum_{n=1}^{\infty}$ — add up the advantage for every possible horizon: 1-step, 2-step, 3-step... all the way out
- w_n — the weight (importance) given to the n -step advantage
- λ (lambda) — a number between 0 and 1 that controls how fast the weights shrink
- λ^{n-1} — means 1-step gets weight 1, 2-step gets weight λ , 3-step gets λ^2 , etc.

What this equation is saying: Instead of committing to one fixed horizon n , GAE takes a weighted average of all horizons at once. The 1-step advantage gets the biggest weight, the 2-step gets slightly less, the 3-step even less, and so on — shrinking by a factor of λ each time. When $\lambda = 0$, only 1-step matters (low variance, more bias). When $\lambda = 1$, all horizons matter equally (low bias, high variance). You tune λ to get the best of both worlds. In practice $N \ll T$ (you do not actually go to infinity — you cut off early), and cutting earlier reduces variance at the cost of a little bias.

Page 14 — The Full PPO Algorithm

1. **Collect data:** Run the current policy π_θ in the environment to gather many episodes.
2. **Fit critic:** Train the value network $\hat{V}_\phi^\pi$ to match the observed summed rewards.
3. **Compute GAE advantage:** Use the n -step formula with weights $w_n \propto \lambda^{n-1}$ to get $\hat{A}_{\text{GAE}}^\pi$ for every action taken.

4. **Update policy:** Run M gradient steps (not just one!) on the full clipped and min surrogate objective from page 13.

The key payoff of all this machinery is Step 4: because PPO has safe clipping, you can squeeze out multiple gradient updates from a single batch of data without the policy flying off the rails. That is something vanilla policy gradients could not do — they had to collect new data after every single update.

The Big Picture Across All Three Pages

Think of it like a student who took one test (collected data under old policy θ). A naive student studies that test to death and overcorrects — completely changing how they think. PPO is the careful student: they use clipping to say “I will only shift my thinking so far from what I believed before the test,” use GAE to get smart credit assignment for each answer, and then review the same test paper multiple times safely — getting more learning out of each exam without going off the rails.

Give me a concrete example of trick 3 to explain it.

Concrete Example: GAE (Trick #3)

Scenario: A robot is learning to walk. At moment t , it steps forward. We want to know: was that step a good action?

Given numbers ($\gamma = 1$, no discounting, to keep it clean):

	Value
Reward at step t	+2 (moved forward)
Reward at step $t + 1$	+1 (another step)
Reward at step $t + 2$	−8 (fell down)
Critic’s guess for $V(s_t)$	3
Critic’s guess for $V(s_{t+1})$	4
Critic’s guess for $V(s_{t+2})$	5
Critic’s guess for $V(s_{t+3})$	0

Step 1 — Compute each n -step advantage

1-step (looks only 1 step ahead):

$$A_1 = r_t + V(s_{t+1}) - V(s_t) = 2 + 4 - 3 = +3$$

“Stepped forward, got +2, next state looks good. Action seems great!”

2-step (looks 2 steps ahead):

$$A_2 = r_t + r_{t+1} + V(s_{t+2}) - V(s_t) = 2 + 1 + 5 - 3 = +5$$

“Two good steps, state ahead looks even better. Looks amazing!” — but the fall is not visible yet.

3-step (looks 3 steps ahead):

$$A_3 = r_t + r_{t+1} + r_{t+2} + V(s_{t+3}) - V(s_t) = 2 + 1 + (-8) + 0 - 3 = -8$$

“Now we can see the robot falls at step $t + 2$. The action looks terrible.”

Step 2 — GAE blends all three with λ -weights

Weights: $w_1 = 1$, $w_2 = \lambda$, $w_3 = \lambda^2$

$$A_{\text{GAE}} = 1 \cdot A_1 + \lambda \cdot A_2 + \lambda^2 \cdot A_3$$

λ value	Meaning	Calculation	Result
$\lambda = 0$	Only 1-step	$1(+3)$	+3 — blind to the fall
$\lambda = 0.5$	Balanced	$1(+3) + 0.5(+5) + 0.25(-8)$	+3.5 — fall partially visible
$\lambda = 0.9$	Looks far ahead	$1(+3) + 0.9(+5) + 0.81(-8)$	+1.02 — fall heavily weighted
$\lambda = 1$	All equal weight	$1(+3) + 1(+5) + 1(-8)$	0 — fall fully cancels the gains

The punchline

The 1-step alone would tell the robot “stepping forward was a great move (+3)!” — because it cannot see the fall that happens two steps later.

The 3-step alone is too noisy — it blames the entire step forward harshly (−8) even though the first two steps were fine.

GAE with $\lambda = 0.9$ says +1.02 — which is honest: the step forward was slightly good, but not as good as it looked short-term, because trouble is coming. That is exactly the kind of precise, blame-accurate signal the actor needs to learn well.

Page 15 — The Big Question: Can We Reuse Old Data?

So far, the actor-critic algorithm only learns from data it just collected. Once you are done with a batch of experience, you throw it away. That is wasteful. Page 15 asks: What if we saved every single experience the agent ever had, and learned from ALL of it — not just the latest batch?

Two key ideas introduced:

1. **Replay Buffer** — a giant diary that stores every (state, action, reward, next-state) experience the agent has ever seen.
2. **Fix the equations** — the original equations assume you only use fresh data. If you use old data, you need to adjust the math.

Page 16 — First Attempt: Plug the Replay Buffer In

They take the standard algorithm and simply say: instead of doing steps on one fresh batch, do them on a small random sample (“minibatch”) pulled from the replay buffer — which contains all past experience.

The 5-step algorithm (first attempt):

1. Collect fresh experience and add it to the replay buffer.
2. Fit the value function $\hat{V}_\phi^\pi$ to reward sums from the buffer.

3. Calculate how good each (state, action) was using the advantage:

$$\hat{A}^\pi(s_i, a_i) = r(s_i, a_i) + \hat{V}_\phi^\pi(s'_i) - \hat{V}_\phi^\pi(s_i)$$

4. Estimate the policy gradient.
5. Update the policy using that gradient.

Steps 2–4 are done on a minibatch randomly sampled from all previous data — that is the key change.

Page 17 — Two Things Are Broken

Once they lay out the algorithm, they realize there is a serious problem. When V^π is fit on ALL data from the replay buffer (which comes from many different past versions of the policy), the value function represents some kind of average over old policies — not the current policy. This is a problem because the value function is supposed to tell you how good the *current* policy is.

Two things broken:

- **The target is wrong.** The value estimate $y = r + \gamma \hat{V}(s')$ uses a $\hat{V}$ trained on a mix of old policies, so it does not correctly represent the current policy's future rewards.
- **The likelihood is wrong.** In the gradient step, you are using the log probability of actions that old versions of the policy took — not actions the current policy would take.

Page 18 — The Fix: Use $Q(s, a)$ Instead of $V(s)$

Core insight: Instead of estimating how good a state is (V), estimate how good a state + action pair is (Q). This matters because $Q(s, a)$ can be computed from any experience, regardless of which policy collected it.

The equation:

$$Q^{\pi_\theta}(s, a) = r(s, a) + \gamma \mathbb{E}_{\substack{s' \sim p(\cdot | s, a) \\ \bar{a}' \sim \pi_\theta(\cdot | s')}} [Q^{\pi_\theta}(s', \bar{a}')]]$$

Symbol table:

Symbol	What it means
$Q^{\pi_\theta}(s, a)$	Q-value: expected total future reward starting in state s , taking action a , then following policy π_θ
$r(s, a)$	Immediate reward for taking action a in state s
γ	Discount factor (0 to 1); sooner rewards are worth more than later ones
$\mathbb{E}$	Expected value — average over many possible outcomes
$s' \sim p(\cdot \mid s, a)$	Next state sampled from the environment's transition function
$\bar{a}' \sim \pi_\theta(\cdot \mid s')$	Next action sampled from the <i>current</i> policy at next state s'
$Q^{\pi_\theta}(s', \bar{a}')$	Q-value of the next state-action pair

The quality of taking action a in state s equals the immediate reward right now, plus the discounted quality of whatever state and action comes next. The “next action” $\bar{a}'$ must be sampled from the *current* policy — not retrieved from the replay buffer — because we want to know how good the current policy is, not some past version. The approach: sample (s, a, s') from the replay buffer (old data is fine), but sample $\bar{a}'$ fresh from the current policy. Then the target becomes $y = r(s, a) + \gamma Q(s', \bar{a}')$, and you train the Q-network toward this target.

Page 19 — Fixing the Value Function: The Updated Algorithm

This page rewrites the algorithm with the fix from page 18 applied.

The key fix in step 3:

- **Old (broken):** Update $\hat{V}_\phi^\pi$ using targets $y_i = r_i + \gamma \hat{V}_\phi^\pi(s'_i)$
- **New (fixed):** Update $\hat{Q}_\phi^\pi$ using targets $y_i = r(s_i, a_i) + \gamma \hat{Q}_\phi^\pi(s'_i, a'_i)$, where $a'_i \sim \pi_\theta(\cdot \mid s'_i)$ — sampled from the *current* policy, not from the replay buffer.

Using a fresh action fixes the broken target problem — the Q-value target now reflects the current policy. The states s_i can still come from the replay buffer; only the next actions need to be fresh.

Page 20 — Two More Refinements

Refinement 1: Drop the baseline, just use Q directly.

Instead of computing the advantage $\hat{A}(s, a) = Q(s, a) - V(s)$, just use $Q(s, a)$ directly. This has higher variance (noisier estimates), but since you now have a huge replay buffer full of data, the extra noise is acceptable — more data cancels it out.

Refinement 2: Sample fresh actions for the gradient step too.

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_i \nabla_\theta \log \pi_\theta(\bar{a}_i \mid s_i) \hat{Q}^\pi(s_i, \bar{a}_i)$$

Symbol table:

Symbol	What it means
$\nabla_{\theta} J(\theta)$	Gradient: the direction to nudge policy parameters θ to improve total reward
$\frac{1}{N} \sum_i$	Average over N sampled examples
$\nabla_{\theta} \log \pi_{\theta}(\bar{a}_i \mid s_i)$	How much the log-probability of action $\bar{a}_i$ changes when you tweak θ
$\bar{a}_i$	Fresh action sampled from current policy π_{θ} at state s_i — NOT the old stored action
$\hat{Q}^{\pi}(s_i, \bar{a}_i)$	Estimated quality of taking fresh action $\bar{a}_i$ in state s_i
$\pi_{\theta}(\bar{a}_i \mid s_i)$	Probability the current policy assigns to action $\bar{a}_i$ given state s_i

To improve the policy, look at a bunch of state-action pairs. For each state s_i from the buffer, sample a fresh action $\bar{a}_i$ from the current policy (not the old stored action). Check how good that fresh action is according to $\hat{Q}$. Then nudge the policy parameters in the direction that makes good actions more likely. Using a fresh action fixes the “wrong likelihood” problem from page 17 — you are now computing log probabilities under the current policy, not some old one. The states s_i can still come from the replay buffer; only the actions need to be fresh.

Page 21 — The Final Clean Algorithm

This is the polished, fixed version of the off-policy actor-critic with a replay buffer.

Final 5-step algorithm:

1. Collect experience $\{s_i, a_i\}$ from the current policy and add it to the replay buffer.
2. Sample a batch $\{s_i, a_i, r_i, s'_i\}$ from the replay buffer.
3. Update $\hat{Q}_{\phi}^{\pi}$ toward targets $y_i = r(s_i, a_i) + \gamma \hat{Q}_{\phi}^{\pi}(s'_i, a'_i)$, where $a'_i \sim \pi_{\theta}(\cdot \mid s'_i)$ — sampled fresh from the current policy.
4. Compute the gradient:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \nabla_{\theta} \log \pi_{\theta}(\bar{a}_i \mid s_i) \hat{Q}^{\pi}(s_i, \bar{a}_i)$$

where $\bar{a}_i \sim \pi_{\theta}(\cdot \mid s_i)$ — freshly sampled from the current policy.

5. Update $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$.

The replay buffer lets the agent learn from every experience it has ever had, making training much more data-efficient. The two key fixes — sampling fresh actions for the Q-target and for the gradient — ensure that even though the *data* is old, the *evaluation* always reflects what the current policy would do.

What is meant by random sample (minibatch)?

The replay buffer is like a giant jar full of thousands of past experiences (state, action, reward, next-state). Instead of learning from every single one of those experiences at once (too slow and memory-heavy), you randomly grab a small handful — say, 64 or 256 — and learn from just those.

That small random handful is the **minibatch**.

“Random” is important here: you do not pick the most recent ones or the best ones — you grab them randomly so the batch is not biased toward any one policy or time period. This also breaks correlations between nearby experiences, which helps the learning stay stable.

$Q(s,a)$ can be computed from any experience. How and why?

First, recall what each one asks:

- $V(s)$: “How good is this state, if I follow my current policy from here?” — depends entirely on the policy.
- $Q(s,a)$: “How good is it to take action a in state s , and then follow my current policy from there?” — the action is already pinned down.

Why Q can use old data:

The target for Q is:

$$y = r(s, a) + \gamma \hat{Q}^\pi(s', \bar{a}')$$

To compute this, you need three things:

Thing needed	Where it comes from	Policy-dependent?
$r(s, a)$	Replay buffer (old data)	No — the environment gave this reward, not the policy
s'	Replay buffer (old data)	No — the environment transitioned to s' , not the policy
$\bar{a}'$	Sampled fresh from current policy	Yes — and you do this fresh

The reward $r(s, a)$ is a fact about the environment — it does not matter which policy decided to take action a . If you step off a cliff in a game, you lose a life whether your old policy did it or your new one. The environment always gives the same reward for the same (s, a) . The next state s' is also determined by the environment, not the policy. So the (s, a, r, s') tuple from old experience is completely valid. The only policy-sensitive piece — what action comes next — is handled by sampling $\bar{a}'$ fresh from the current policy at s' .

Why V cannot use old data:

$V(s)$ has no action pinned. It averages over whatever actions the policy would take. So if your replay buffer has data from 10 different old policies, $V(s)$ ends up representing some confused blend of all of them — not your current policy. There is no way to “correct” for this without knowing exactly which old policy produced each experience.

In short: Q separates the policy-independent part (r, s') from the policy-dependent part $(\bar{a}')$. Old data handles the first part; the current policy handles the second. V has no such clean separation.